

# ABDULLA SAJAD

+91 9207761088 | [abdullasajad01@gmail.com](mailto:abdullasajad01@gmail.com) | [neuralbroker.github.io](https://neuralbroker.github.io) | [github.com/neuralbroker](https://github.com/neuralbroker)

## EDUCATION

**Sree Narayana Guru College of Engineering and Technology**

*B.Tech in Computer Science and Engineering*

Kerala, India

Sep. 2022 – May 2026

## TECHNICAL SKILLS

**AI / LLM:** PyTorch, Fine-tuning (SFT, GRPO, DPO, LoRA), LangChain, LangGraph, RAG, FAISS, Hugging Face, GGUF/llama.cpp, OpenAI API

**Languages:** Python, TypeScript, Java, SQL

**Backend:** FastAPI, REST APIs, WebSockets, Celery, PostgreSQL, Redis, SQLAlchemy

**Infra & MLOps:** Docker, GitHub Actions, MLflow, AWS (EC2), Git, Linux, Supabase, Vercel

**Frontend:** React, Next.js, Tailwind CSS

## PROJECTS

**BlitzKode** | *Python, PyTorch, Qwen2.5-Coder-3B, SFT, GRPO, DPO, LoRA, GGUF, Hugging Face* | [HF](#)

- Implemented SFT → GRPO (the RL alignment technique in DeepSeek-R1) → DPO on Qwen2.5-Coder-3B using PyTorch and Hugging Face; cut training VRAM from ~24GB to 8GB via LoRA, gradient checkpointing, and bfloat16 mixed precision.
- Quantised to GGUF for CPU-only inference; published as *neuralbroker/blitzkode* on Hugging Face with **510+ downloads** and fully reproducible training configs.

**EsenceLab** | *Next.js, TypeScript, FastAPI, PostgreSQL, Supabase, Gemini AI, Docker* | [Live](#)

- Led a cross-functional 4-engineer team across frontend, backend, and AI components to build a campus recruitment platform – Gemini-powered resume parsing, skill-gap analysis, and candidate ranking cut manual shortlisting by ~60% across 200+ resumes; **won 2nd Prize at Innovision (30+ teams)**.
- Designed async Celery + FastAPI backend with JWT + Supabase RLS auth and GitHub Actions CI/CD; pre-launch audit eliminated hardcoded credentials and injection vectors. Live, zero incidents.

**SentinelML** | *Python, FastAPI, PyTorch, XGBoost, LightGBM, scikit-learn, Redis, PostgreSQL, MLflow, Docker*

- Built a fraud detection pipeline on imbalanced transaction data: 20+ engineered features, SMOTE oversampling, stacked XGBoost + LightGBM ensemble – ROC-AUC 0.94 on held-out test set.
- Served predictions via FastAPI with Redis caching for low-latency inference; MLflow experiment registry tracks every run for full reproducibility, with versioned model rollback in under 60 seconds. Open-source.

**Lexora AI** | *Python, FastAPI, LangChain, FAISS, OpenAI API*

- Built a RAG pipeline over private document collections using LangChain + FAISS; chunked retrieval with streaming responses achieves sub-2s latency on corpora up to 100k tokens, grounding answers in source documents to eliminate hallucination.

## EXPERIENCE

**Project Lead**

Aug 2024 – May 2026

*SNGCET*

Kerala, India

- Led a cross-functional 4-engineer team across two projects (SpeakSwap, EsenceLab) – sprint planning, architecture, PR reviews, and GitHub Actions CI/CD; shipped both to production on schedule with zero regressions.

**Freelance Full Stack Developer**

Jun 2023 – 2025

*Self-Employed*

Remote

- Delivered two production web applications independently – designed data models, built REST APIs, and managed deployments end-to-end for paying clients.

## ACHIEVEMENTS

**2nd Prize – Innovision Project Competition** | EsenceLab, SNGCET (30+ competing teams)

**510+ downloads** | *neuralbroker/blitzkode* on Hugging Face

**Seminar – Neurosymbolic AI in Military Systems** | Explainability and safety in mission-critical AI